

```

In [2]: import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms, models
from PIL import Image
import pandas as pd
import os
import numpy as np
from sklearn.model_selection import train_test_split
from tqdm import tqdm

def preprocess_image(image):
    image = image.convert("L")
    image = np.array(image)
    local_mean = np.mean(image)
    thresh_image = np.where(image > (local_mean - 2), 255, 0)
    return Image.fromarray(thresh_image.astype(np.uint8))

class GlomeruliDataset(Dataset):
    def __init__(self, csv_file, img_dir, transform=None):
        self.annotations = pd.read_csv(csv_file)
        self.img_dir = img_dir
        self.transform = transform

    def __len__(self):
        return len(self.annotations)

    def __getitem__(self, index):
        img_path = os.path.join(self.img_dir, self.annotations.iloc[index, 0])
        image = Image.open(img_path)
        image = preprocess_image(image)
        image = image.convert("RGB")
        y_label = torch.tensor(int(self.annotations.iloc[index, 1]))

        if self.transform:
            image = self.transform(image)
        return image, y_label

class GlomeruliClassifier(nn.Module):
    def __init__(self, base_model, num_classes=2):
        super(GlomeruliClassifier, self).__init__()
        self.base_model = base_model
        num_features = 2048
        self.classifier = nn.Sequential(
            nn.Linear(num_features, 512),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(512, num_classes),
        )

```

```

def forward(self, x):
    features = self.base_model(x)
    return self.classifier(features)

def train_model(
    model, train_loader, val_loader, criterion, optimizer, scheduler, num_epochs, d
):
    best_loss = float("inf")
    best_accuracy = 0.0

    for epoch in range(num_epochs):
        # Training phase
        model.train()
        running_loss = 0.0
        progress_bar = tqdm(train_loader, desc=f"Epoch {epoch+1}/{num_epochs}")

        for images, labels in progress_bar:
            images, labels = images.to(device), labels.to(device)
            optimizer.zero_grad()
            outputs = model(images)
            loss = criterion(outputs, labels)
            loss.backward()
            optimizer.step()
            running_loss += loss.item()

            progress_bar.set_postfix(
                {
                    "Loss": f"{running_loss/(progress_bar.n+1):.4f}",
                    "LR": f'{optimizer.param_groups[0]["lr"]:.6f}',
                }
            )

        epoch_loss = running_loss / len(train_loader)

        # Validation phase
        model.eval()
        val_loss = 0.0
        correct = 0
        total = 0

        with torch.no_grad():
            for images, labels in val_loader:
                images, labels = images.to(device), labels.to(device)
                outputs = model(images)
                loss = criterion(outputs, labels)
                val_loss += loss.item()
                _, predicted = outputs.max(1)
                total += labels.size(0)
                correct += predicted.eq(labels).sum().item()

        val_loss /= len(val_loader)
        val_accuracy = 100.0 * correct / total

        print(f"\nEpoch {epoch+1}/{num_epochs}")
        print(f"Average Loss: {epoch_loss:.4f}")

```

```

print(f"Validation Loss: {val_loss:.4f}")
print(f"Validation Accuracy: {val_accuracy:.2f}%")

if val_loss < best_loss:
    best_loss = val_loss
    torch.save(model.state_dict(), "Final_model.pth")
    print(f"Saved new best model with validation loss: {best_loss:.4f}")

scheduler.step(val_loss)
print("\n")

if __name__ == "__main__":
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

    transform = transforms.Compose(
        [
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        ]
    )

    dataset = GlomeruliDataset(
        csv_file="./train_labels.csv",
        img_dir="./ResizedTrainingSet",
        transform=transform,
    )

    labels = [int(dataset.annotations.iloc[i, 1]) for i in range(len(dataset))]
    class_counts = np.bincount(labels)
    total_samples = len(labels)
    class_weights = total_samples / (len(class_counts) * class_counts)
    class_weights = torch.FloatTensor(class_weights).to(device)
    print("Class distribution:", class_counts)
    print("Class weights:", class_weights.cpu().numpy())

    train_set, val_set = train_test_split(dataset, test_size=0.2, random_state=42)
    train_loader = DataLoader(train_set, batch_size=32, shuffle=True, num_workers=2)
    val_loader = DataLoader(val_set, batch_size=32, shuffle=False, num_workers=2)

    model_path = "./inception_v3_google-0cc3c7bd.pth"
    base_model = models.inception_v3(pretrained=False)
    base_model.load_state_dict(torch.load(model_path))
    base_model.aux_logits = False
    base_model.fc = nn.Identity()
    model = GlomeruliClassifier(base_model).to(device)

    criterion = nn.CrossEntropyLoss(weight=class_weights)
    optimizer = optim.Adam(model.parameters(), lr=0.001)
    scheduler = optim.lr_scheduler.ReduceLROnPlateau(
        optimizer, mode="min", factor=0.1, patience=3, verbose=True
    )

    train_model(

```

```

    model,
    train_loader,
    val_loader,
    criterion,
    optimizer,
    scheduler,
    num_epochs=20,
    device=device,
)

```

Using device: cuda

Class distribution: [3763 843]

Class weights: [0.6120117 2.7319098]

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:208: UserWarning: The parameter 'pretrained' is deprecated since 0.13 and may be removed in the future, please use 'weights' instead.

warnings.warn(

/usr/local/lib/python3.10/dist-packages/torchvision/models/_utils.py:223: UserWarning: Arguments other than a weight enum or `None` for 'weights' are deprecated since 0.13 and may be removed in the future. The current behavior is equivalent to passing `weights=None`.

warnings.warn(msg)

/usr/local/lib/python3.10/dist-packages/torchvision/models/inception.py:43: FutureWarning: The default weight initialization of inception_v3 will be changed in future releases of torchvision. If you wish to keep the old behavior (which leads to long initialization times due to scipy/scipy#11299), please set init_weights=True.

warnings.warn(

<ipython-input-2-86d37c2338d1>:155: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

base_model.load_state_dict(torch.load(model_path))

/usr/local/lib/python3.10/dist-packages/torch/optim/lr_scheduler.py:60: UserWarning: The verbose parameter is deprecated. Please use get_last_lr() to access the learning rate.

warnings.warn(

Epoch 1/20: 100%|██████████| 116/116 [00:50<00:00, 2.28it/s, Loss=0.3494, LR=0.001000]

Epoch 1/20

Average Loss: 0.3494

Validation Loss: 0.1550

Validation Accuracy: 93.17%

Saved new best model with validation loss: 0.1550

Epoch 2/20: 100%|██████████| 116/116 [00:52<00:00, 2.22it/s, Loss=0.2473, LR=0.001000]

Epoch 2/20
Average Loss: 0.2473
Validation Loss: 0.1519
Validation Accuracy: 93.82%
Saved new best model with validation loss: 0.1519

Epoch 3/20: 100%|██████████| 116/116 [00:53<00:00, 2.18it/s, Loss=0.1861, LR=0.001000]

Epoch 3/20
Average Loss: 0.1861
Validation Loss: 0.3580
Validation Accuracy: 80.69%

Epoch 4/20: 100%|██████████| 116/116 [00:54<00:00, 2.15it/s, Loss=0.1221, LR=0.001000]

Epoch 4/20
Average Loss: 0.1221
Validation Loss: 0.2450
Validation Accuracy: 93.28%

Epoch 5/20: 100%|██████████| 116/116 [00:54<00:00, 2.12it/s, Loss=0.1136, LR=0.001000]

Epoch 5/20
Average Loss: 0.1136
Validation Loss: 0.2634
Validation Accuracy: 94.58%

Epoch 6/20: 100%|██████████| 116/116 [00:55<00:00, 2.11it/s, Loss=0.1631, LR=0.001000]

Epoch 6/20
Average Loss: 0.1631
Validation Loss: 0.2793
Validation Accuracy: 87.96%

Epoch 7/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0729, LR=0.000100]

Epoch 7/20
Average Loss: 0.0729
Validation Loss: 0.1180
Validation Accuracy: 96.10%
Saved new best model with validation loss: 0.1180

Epoch 8/20: 100%|██████████| 116/116 [00:55<00:00, 2.10it/s, Loss=0.0490, LR=0.000100]

Epoch 8/20
Average Loss: 0.0490
Validation Loss: 0.1560
Validation Accuracy: 96.20%

Epoch 9/20: 100%|██████████| 116/116 [00:55<00:00, 2.10it/s, Loss=0.0348, LR=0.000100]

Epoch 9/20
Average Loss: 0.0348
Validation Loss: 0.1573
Validation Accuracy: 96.10%

Epoch 10/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0263, LR=0.000100]

Epoch 10/20
Average Loss: 0.0263
Validation Loss: 0.1559
Validation Accuracy: 96.10%

Epoch 11/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0291, LR=0.000100]

Epoch 11/20
Average Loss: 0.0291
Validation Loss: 0.1931
Validation Accuracy: 96.20%

Epoch 12/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0138, LR=0.000010]

Epoch 12/20
Average Loss: 0.0138
Validation Loss: 0.1798
Validation Accuracy: 95.44%

Epoch 13/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0096, LR=0.000010]

Epoch 13/20
Average Loss: 0.0096
Validation Loss: 0.2013
Validation Accuracy: 95.66%

Epoch 14/20: 100%|██████████| 116/116 [00:55<00:00, 2.08it/s, Loss=0.0193, LR=0.000010]

Epoch 14/20
Average Loss: 0.0193
Validation Loss: 0.2241
Validation Accuracy: 96.42%

Epoch 15/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0118, LR=0.000010]

Epoch 15/20

Average Loss: 0.0118

Validation Loss: 0.1762

Validation Accuracy: 95.66%

Epoch 16/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0098, LR=0.000001]

Epoch 16/20

Average Loss: 0.0098

Validation Loss: 0.1782

Validation Accuracy: 95.99%

Epoch 17/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0178, LR=0.000001]

Epoch 17/20

Average Loss: 0.0178

Validation Loss: 0.1903

Validation Accuracy: 96.10%

Epoch 18/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0131, LR=0.000001]

Epoch 18/20

Average Loss: 0.0131

Validation Loss: 0.1878

Validation Accuracy: 95.66%

Epoch 19/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0125, LR=0.000001]

Epoch 19/20

Average Loss: 0.0125

Validation Loss: 0.1884

Validation Accuracy: 95.77%

Epoch 20/20: 100%|██████████| 116/116 [00:55<00:00, 2.09it/s, Loss=0.0111, LR=0.000000]

Epoch 20/20

Average Loss: 0.0111

Validation Loss: 0.1644

Validation Accuracy: 95.99%

```
In [5]: import torch
import torch.nn as nn
from torch.utils.data import Dataset, DataLoader
from torchvision import transforms, models
from PIL import Image
```

```

import pandas as pd
import os
import numpy as np
from tqdm import tqdm
from sklearn.metrics import classification_report, confusion_matrix

def preprocess_image(image):
    image = image.convert("L")
    image = np.array(image)
    local_mean = np.mean(image)
    thresh_image = np.where(image > (local_mean - 2), 255, 0)
    return Image.fromarray(thresh_image.astype(np.uint8))

class GlomeruliDataset(Dataset):
    def __init__(self, csv_file, img_dir, transform=None):
        self.annotations = pd.read_csv(csv_file)
        self.img_dir = img_dir
        self.transform = transform

    def __len__(self):
        return len(self.annotations)

    def __getitem__(self, index):
        img_path = os.path.join(self.img_dir, self.annotations.iloc[index, 0])
        image = Image.open(img_path)
        image = preprocess_image(image)
        image = image.convert("RGB")
        y_label = torch.tensor(int(self.annotations.iloc[index, 1]))
        if self.transform:
            image = self.transform(image)
        return image, y_label

class GlomeruliClassifier(nn.Module):
    def __init__(self, base_model, num_classes=2):
        super(GlomeruliClassifier, self).__init__()
        self.base_model = base_model
        num_features = 2048
        self.classifier = nn.Sequential(
            nn.Linear(num_features, 512),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(512, num_classes),
        )

    def forward(self, x):
        features = self.base_model(x)
        return self.classifier(features)

def evaluate_model(model, test_loader, device):
    model.eval()
    correct = 0
    total = 0

```



```

predictions = []
true_labels = []

with torch.no_grad():
    for images, labels in tqdm(test_loader):
        images, labels = images.to(device), labels.to(device)
        outputs = model(images)
        _, predicted = outputs.max(1)

        total += labels.size(0)
        correct += predicted.eq(labels).sum().item()

        predictions.extend(predicted.cpu().numpy())
        true_labels.extend(labels.cpu().numpy())

accuracy = 100.0 * correct / total
return accuracy, predictions, true_labels

if __name__ == "__main__":
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

    # Data transformations
    transform = transforms.Compose(
        [
            transforms.RandomHorizontalFlip(),
            transforms.ToTensor(),
            transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
        ]
    )

    # Test data
    test_dataset = GlomeruliDataset(
        csv_file="./test_labels.csv", # Path to your test labels CSV
        img_dir="./ResizedTestSet", # Path to your test images directory
        transform=transform,
    )
    test_loader = DataLoader(test_dataset, batch_size=32, shuffle=False, num_workers=0)

    # Create a models directory in your project
    model_path = "./inception_v3_google-0cc3c7bd.pth"

    # Load the model with custom weights using safe loading
    base_model = models.inception_v3(weights=None, init_weights=True)
    base_model.load_state_dict(torch.load(model_path, weights_only=True))
    base_model.fc = nn.Identity()
    model = GlomeruliClassifier(base_model).to(device)

    # Load the trained model
    model.load_state_dict(torch.load("./Final_model.pth"))

    # Evaluate the model
    accuracy, predictions, true_labels = evaluate_model(model, test_loader, device)

    print(f"\nTest Accuracy: {accuracy:.2f}%")

```

```
print("\nClassification Report:")
print(classification_report(true_labels, predictions))
print("\nConfusion Matrix:")
print(confusion_matrix(true_labels, predictions))
```

Using device: cuda

<ipython-input-5-9b37f231db4e>:112: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```
model.load_state_dict(torch.load("/kaggle/working/Final_model.pth"))
100%|██████████| 36/36 [00:06<00:00, 5.55it/s]
```

Test Accuracy: 95.75%

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.95	0.97	941
1	0.83	0.97	0.89	211
accuracy			0.96	1152
macro avg	0.91	0.96	0.93	1152
weighted avg	0.96	0.96	0.96	1152

Confusion Matrix:

```
[[898  43]
 [  6 205]]
```

```
In [3]: import torch
import torch.nn as nn
from torchvision import transforms, models
from PIL import Image
import os
import numpy as np
import csv
from tqdm import tqdm

def preprocess_image(image):
    image = image.convert("L")
    image = np.array(image)
    local_mean = np.mean(image)
    thresh_image = np.where(image > (local_mean - 2), 255, 0)
    return Image.fromarray(thresh_image.astype(np.uint8))

class GlomeruliClassifier(nn.Module):
    def __init__(self, base_model, num_classes=2):
        super(GlomeruliClassifier, self).__init__()
        self.base_model = base_model
```

```

        num_features = 2048
        self.classifier = nn.Sequential(
            nn.Linear(num_features, 512),
            nn.ReLU(),
            nn.Dropout(0.5),
            nn.Linear(512, num_classes),
        )

    def forward(self, x):
        features = self.base_model(x)
        return self.classifier(features)

def predict_images(model, image_folder, device):
    model.eval()
    predictions = []
    transform = transforms.Compose([
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    with torch.no_grad():
        for image_name in tqdm(os.listdir(image_folder)):
            if image_name.lower().endswith(('.png', '.jpg', '.jpeg')):
                img_path = os.path.join(image_folder, image_name)
                image = Image.open(img_path)
                image = preprocess_image(image)
                image = image.convert("RGB")
                image = transform(image).unsqueeze(0).to(device)

                outputs = model(image)
                _, predicted = outputs.max(1)
                predictions.append((image_name, predicted.item()))

    return predictions

def save_predictions_to_csv(predictions, output_file):
    with open(output_file, 'w', newline='') as csvfile:
        writer = csv.writer(csvfile)
        writer.writerow(['Image Name', 'Prediction'])
        writer.writerows(predictions)

if __name__ == "__main__":
    device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
    print(f"Using device: {device}")

    # Load the model
    model_path = "./inception_v3_google-0cc3c7bd.pth"
    base_model = models.inception_v3(weights=None, init_weights=True)
    base_model.load_state_dict(torch.load(model_path, weights_only=True))
    base_model.fc = nn.Identity()
    model = GlomeruliClassifier(base_model).to(device)

    # Load the trained model weights
    model.load_state_dict(torch.load("./Final_model.pth"))

```

```

# Specify the folder containing the images to be classified
image_folder = "./ResizedTestSet"

# Predict classifications for all images in the folder
predictions = predict_images(model, image_folder, device)

# Save predictions to CSV
output_file = "evaluation.csv"
save_predictions_to_csv(predictions, output_file)

print(f"Predictions saved to {output_file}")

```

Using device: cuda

<ipython-input-3-5b22aa7a7c31>:75: FutureWarning: You are using `torch.load` with `weights_only=False` (the current default value), which uses the default pickle module implicitly. It is possible to construct malicious pickle data which will execute arbitrary code during unpickling (See <https://github.com/pytorch/pytorch/blob/main/SECURITY.md#untrusted-models> for more details). In a future release, the default value for `weights_only` will be flipped to `True`. This limits the functions that could be executed during unpickling. Arbitrary objects will no longer be allowed to be loaded via this mode unless they are explicitly allowlisted by the user via `torch.serialization.add_safe_globals`. We recommend you start setting `weights_only=True` for any use case where you don't have full control of the loaded file. Please open an issue on GitHub for any issues related to this experimental feature.

```
model.load_state_dict(torch.load("/kaggle/working/Final_model.pth"))
```

```
100%|██████████| 1152/1152 [00:30<00:00, 37.65it/s]
```

Predictions saved to evaluation.csv